Held-out policy evaluation

Points: individual episodes | diamonds: mean | bars: approximate 95% confidence interval

Episodes per policy: 200, 200, 200 | matched initial-state seeds | one trained checkpoint

Intervals reflect evaluation episodes; they do not measure variation across independent training runs.